

Car Analysis

October 22, 2019

1 Cars features and price dataset

This is an analysis of the Car features and price dataset from Kaggle. My aims for the dataset are the following - Clean the dataset - Conduct Univariate and Bivariate analysis to understand what factors have an effect on pricing and fuel economy - Use Machine Learning algorithms to find - Predictive features for Fuel Economy - Regression model to predict MSRP

1.1 Import the necessary libraries and data file and perform initial data assessment

```
In [1]: # Import libraries necessary for this project
import numpy as np
import pandas as pd
from time import time
from IPython.display import display # Allows the use of display() for DataFrames
import matplotlib.pyplot as plt
import pandas_profiling
import seaborn as sns
# Pretty display for notebooks
%matplotlib inline
import visuals as vs
```

```
# Load the Census dataset
data = pd.read_csv("data.csv")
```

```
In [2]: #create a profile report
data.profile_report(style={'full_width':True})
```

<IPython.core.display.HTML object>

Out[2]:

I really like using the profiling function, because it gives a lot of information in one like of code, and can help with the data wrangling and univariate analysis

Based on the report we can see - We have missing values in the dataset - We have duplicated values - Highway MPG and City MPG are highly correlated

Based on this information, we can begin the data wrangling process

1.2 Data Wrangling

In this step I will clean up the data 1. Deal with the null values 2. Deal with the outliers

```
In [3]: #Check if the dataset has any null values
        data.isnull().sum()
```

```
Out[3]: Make                0
        Model              0
        Year              0
        Engine_Fuel_Type    3
        Engine_HP          69
        Engine_Cylinders    30
        Transmission_Type   0
        Driven_Wheels       0
        Number_of_Doors     6
        Market_Category     3742
        Vehicle_Size        0
        Vehicle_Style       0
        highway_MPG         0
        city_mpg            0
        Popularity          0
        MSRP                0
        dtype: int64
```

```
In [4]: #Check if there is anything specific that stands out about the null values in the Mark
        data[data['Market_Category'].isnull()]
```

```
Out[4]:
```

	Make	Model	Year	Engine_Fuel_Type	Engine_HP	\
87	Nissan	200SX	1996	regular unleaded	115.0	
88	Nissan	200SX	1996	regular unleaded	115.0	
91	Nissan	200SX	1997	regular unleaded	115.0	
92	Nissan	200SX	1997	regular unleaded	115.0	
93	Nissan	200SX	1998	regular unleaded	115.0	
94	Nissan	200SX	1998	regular unleaded	115.0	
203	Chrysler	300	2015	regular unleaded	300.0	
204	Chrysler	300	2015	regular unleaded	292.0	
205	Chrysler	300	2015	regular unleaded	292.0	
206	Chrysler	300	2015	regular unleaded	292.0	
209	Chrysler	300	2015	regular unleaded	292.0	
210	Chrysler	300	2015	regular unleaded	292.0	
211	Chrysler	300	2016	regular unleaded	300.0	
213	Chrysler	300	2016	regular unleaded	292.0	
214	Chrysler	300	2016	regular unleaded	292.0	
215	Chrysler	300	2016	regular unleaded	292.0	
216	Chrysler	300	2016	regular unleaded	292.0	
219	Chrysler	300	2016	regular unleaded	292.0	
220	Chrysler	300	2016	regular unleaded	292.0	
221	Chrysler	300	2016	regular unleaded	300.0	

222	Chrysler	300	2016	regular unleaded	292.0
223	Chrysler	300	2017	regular unleaded	292.0
224	Chrysler	300	2017	regular unleaded	292.0
225	Chrysler	300	2017	regular unleaded	300.0
228	Chrysler	300	2017	regular unleaded	300.0
229	Chrysler	300	2017	regular unleaded	292.0
231	Chrysler	300	2017	regular unleaded	292.0
360	Mazda	3	2015	regular unleaded	155.0
361	Mazda	3	2015	regular unleaded	155.0
362	Mazda	3	2015	regular unleaded	155.0
...
11686	Suzuki	XL-7	2006	regular unleaded	185.0
11687	Suzuki	XL-7	2006	regular unleaded	185.0
11744	Nissan	Xterra	2013	regular unleaded	261.0
11745	Nissan	Xterra	2013	regular unleaded	261.0
11746	Nissan	Xterra	2013	regular unleaded	261.0
11747	Nissan	Xterra	2013	regular unleaded	261.0
11748	Nissan	Xterra	2013	regular unleaded	261.0
11749	Nissan	Xterra	2013	regular unleaded	261.0
11750	Nissan	Xterra	2013	regular unleaded	261.0
11751	Nissan	Xterra	2014	regular unleaded	261.0
11752	Nissan	Xterra	2014	regular unleaded	261.0
11753	Nissan	Xterra	2014	regular unleaded	261.0
11754	Nissan	Xterra	2014	regular unleaded	261.0
11755	Nissan	Xterra	2014	regular unleaded	261.0
11756	Nissan	Xterra	2014	regular unleaded	261.0
11757	Nissan	Xterra	2014	regular unleaded	261.0
11758	Nissan	Xterra	2015	regular unleaded	261.0
11759	Nissan	Xterra	2015	regular unleaded	261.0
11760	Nissan	Xterra	2015	regular unleaded	261.0
11761	Nissan	Xterra	2015	regular unleaded	261.0
11762	Nissan	Xterra	2015	regular unleaded	261.0
11763	Nissan	Xterra	2015	regular unleaded	261.0
11764	Nissan	Xterra	2015	regular unleaded	261.0
11792	Subaru	XT	1991	regular unleaded	97.0
11793	Subaru	XT	1991	regular unleaded	145.0
11794	Subaru	XT	1991	regular unleaded	145.0
11809	Toyota	Yaris iA	2017	regular unleaded	106.0
11810	Toyota	Yaris iA	2017	regular unleaded	106.0
11867	GMC	Yukon	2015	premium unleaded (recommended)	420.0
11868	GMC	Yukon	2015	premium unleaded (recommended)	420.0

	Engine_Cylinders	Transmission_Type	Driven_Wheels	Number_of_Doors	\
87	4.0	MANUAL	front wheel drive	2.0	
88	4.0	MANUAL	front wheel drive	2.0	
91	4.0	MANUAL	front wheel drive	2.0	
92	4.0	MANUAL	front wheel drive	2.0	
93	4.0	MANUAL	front wheel drive	2.0	

94	4.0	MANUAL	front wheel drive	2.0
203	6.0	AUTOMATIC	all wheel drive	4.0
204	6.0	AUTOMATIC	rear wheel drive	4.0
205	6.0	AUTOMATIC	rear wheel drive	4.0
206	6.0	AUTOMATIC	all wheel drive	4.0
209	6.0	AUTOMATIC	all wheel drive	4.0
210	6.0	AUTOMATIC	all wheel drive	4.0
211	6.0	AUTOMATIC	all wheel drive	4.0
213	6.0	AUTOMATIC	all wheel drive	4.0
214	6.0	AUTOMATIC	rear wheel drive	4.0
215	6.0	AUTOMATIC	all wheel drive	4.0
216	6.0	AUTOMATIC	all wheel drive	4.0
219	6.0	AUTOMATIC	rear wheel drive	4.0
220	6.0	AUTOMATIC	rear wheel drive	4.0
221	6.0	AUTOMATIC	all wheel drive	4.0
222	6.0	AUTOMATIC	all wheel drive	4.0
223	6.0	AUTOMATIC	all wheel drive	4.0
224	6.0	AUTOMATIC	all wheel drive	4.0
225	6.0	AUTOMATIC	all wheel drive	4.0
228	6.0	AUTOMATIC	all wheel drive	4.0
229	6.0	AUTOMATIC	rear wheel drive	4.0
231	6.0	AUTOMATIC	all wheel drive	4.0
360	4.0	AUTOMATIC	front wheel drive	4.0
361	4.0	MANUAL	front wheel drive	4.0
362	4.0	MANUAL	front wheel drive	4.0
...
11686	6.0	AUTOMATIC	rear wheel drive	4.0
11687	6.0	AUTOMATIC	rear wheel drive	4.0
11744	6.0	AUTOMATIC	four wheel drive	4.0
11745	6.0	MANUAL	four wheel drive	4.0
11746	6.0	MANUAL	four wheel drive	4.0
11747	6.0	AUTOMATIC	four wheel drive	4.0
11748	6.0	AUTOMATIC	rear wheel drive	4.0
11749	6.0	AUTOMATIC	four wheel drive	4.0
11750	6.0	AUTOMATIC	rear wheel drive	4.0
11751	6.0	AUTOMATIC	four wheel drive	4.0
11752	6.0	AUTOMATIC	rear wheel drive	4.0
11753	6.0	AUTOMATIC	four wheel drive	4.0
11754	6.0	AUTOMATIC	four wheel drive	4.0
11755	6.0	MANUAL	four wheel drive	4.0
11756	6.0	AUTOMATIC	rear wheel drive	4.0
11757	6.0	MANUAL	four wheel drive	4.0
11758	6.0	MANUAL	four wheel drive	4.0
11759	6.0	AUTOMATIC	four wheel drive	4.0
11760	6.0	AUTOMATIC	four wheel drive	4.0
11761	6.0	MANUAL	four wheel drive	4.0
11762	6.0	AUTOMATIC	rear wheel drive	4.0
11763	6.0	AUTOMATIC	four wheel drive	4.0

11764	6.0	AUTOMATIC	rear wheel drive	4.0
11792	4.0	MANUAL	front wheel drive	2.0
11793	6.0	AUTOMATIC	front wheel drive	2.0
11794	6.0	MANUAL	all wheel drive	2.0
11809	4.0	MANUAL	front wheel drive	4.0
11810	4.0	AUTOMATIC	front wheel drive	4.0
11867	8.0	AUTOMATIC	rear wheel drive	4.0
11868	8.0	AUTOMATIC	four wheel drive	4.0

	Market_Category	Vehicle_Size	Vehicle_Style	highway_MPG	city_mpg \
87	NaN	Compact	Coupe	36	26
88	NaN	Compact	Coupe	36	26
91	NaN	Compact	Coupe	35	25
92	NaN	Compact	Coupe	35	25
93	NaN	Compact	Coupe	35	25
94	NaN	Compact	Coupe	35	25
203	NaN	Large	Sedan	27	18
204	NaN	Large	Sedan	31	19
205	NaN	Large	Sedan	31	19
206	NaN	Large	Sedan	27	18
209	NaN	Large	Sedan	27	18
210	NaN	Large	Sedan	27	18
211	NaN	Large	Sedan	27	18
213	NaN	Large	Sedan	27	18
214	NaN	Large	Sedan	31	19
215	NaN	Large	Sedan	27	18
216	NaN	Large	Sedan	27	18
219	NaN	Large	Sedan	31	19
220	NaN	Large	Sedan	31	19
221	NaN	Large	Sedan	27	18
222	NaN	Large	Sedan	27	18
223	NaN	Large	Sedan	27	18
224	NaN	Large	Sedan	27	18
225	NaN	Large	Sedan	27	18
228	NaN	Large	Sedan	27	18
229	NaN	Large	Sedan	30	19
231	NaN	Large	Sedan	27	18
360	NaN	Compact	Sedan	41	30
361	NaN	Compact	Sedan	41	29
362	NaN	Compact	Sedan	41	29
...
11686	NaN	Midsize	4dr SUV	21	16
11687	NaN	Midsize	4dr SUV	21	16
11744	NaN	Midsize	4dr SUV	20	15
11745	NaN	Midsize	4dr SUV	20	16
11746	NaN	Midsize	4dr SUV	20	16
11747	NaN	Midsize	4dr SUV	20	15
11748	NaN	Midsize	4dr SUV	22	16

11749	NaN	Midsize	4dr SUV	20	15
11750	NaN	Midsize	4dr SUV	22	16
11751	NaN	Midsize	4dr SUV	20	15
11752	NaN	Midsize	4dr SUV	22	16
11753	NaN	Midsize	4dr SUV	20	15
11754	NaN	Midsize	4dr SUV	20	15
11755	NaN	Midsize	4dr SUV	20	16
11756	NaN	Midsize	4dr SUV	22	16
11757	NaN	Midsize	4dr SUV	20	16
11758	NaN	Midsize	4dr SUV	20	15
11759	NaN	Midsize	4dr SUV	20	15
11760	NaN	Midsize	4dr SUV	20	15
11761	NaN	Midsize	4dr SUV	20	15
11762	NaN	Midsize	4dr SUV	22	16
11763	NaN	Midsize	4dr SUV	20	15
11764	NaN	Midsize	4dr SUV	22	16
11792	NaN	Compact	Coupe	29	22
11793	NaN	Compact	Coupe	26	18
11794	NaN	Compact	Coupe	23	16
11809	NaN	Compact	Sedan	39	30
11810	NaN	Compact	Sedan	40	32
11867	NaN	Large	4dr SUV	21	15
11868	NaN	Large	4dr SUV	21	14

	Popularity	MSRP
87	2009	2000
88	2009	2000
91	2009	2000
92	2009	2000
93	2009	2000
94	2009	2000
203	1013	37570
204	1013	31695
205	1013	38070
206	1013	44895
209	1013	34195
210	1013	40570
211	1013	38095
213	1013	45190
214	1013	32260
215	1013	37755
216	1013	41055
219	1013	38555
220	1013	35255
221	1013	38590
222	1013	34760
223	1013	41135
224	1013	45270

225	1013	38670
228	1013	38175
229	1013	32340
231	1013	34840
360	586	23795
361	586	19595
362	586	18445
...
11686	481	25499
11687	481	21999
11744	2009	26900
11745	2009	29440
11746	2009	25850
11747	2009	30490
11748	2009	24850
11749	2009	24990
11750	2009	22940
11751	2009	31370
11752	2009	25300
11753	2009	27350
11754	2009	25440
11755	2009	26300
11756	2009	23390
11757	2009	30320
11758	2009	26670
11759	2009	27720
11760	2009	25710
11761	2009	30590
11762	2009	23660
11763	2009	31640
11764	2009	25670
11792	640	2000
11793	640	2000
11794	640	2000
11809	2031	15950
11810	2031	17050
11867	549	64520
11868	549	67520

[3742 rows x 16 columns]

From the table above it seems that there is nothing specific about the Nan's in the Market Category Column. Since, market category is not an independent characteristic (depends on other factors, like the make, model, style etc). Therefore, for downstream analysis, we can drop this column.

```
In [5]: #drop the market category column
data = data.drop(['Market_Category'], axis=1)
data.isnull().sum()
```

```

Out [5]: Make           0
        Model          0
        Year           0
        Engine_Fuel_Type 3
        Engine_HP       69
        Engine_Cylinders 30
        Transmission_Type 0
        Driven_Wheels    0
        Number_of_Doors  6
        Vehicle_Size     0
        Vehicle_Style     0
        highway_MPG      0
        city_mpg         0
        Popularity       0
        MSRP             0
        dtype: int64

```

1.2.1 Null Values

The two major sources of null values are horsepower and engine cylinders. We will look at both of them and see if we can add the missing values to these or would we have to drop them

Horsepower

```

In [6]: #Create a dataframe to further investigate the null values of Engine HP
        df_temp = data[data['Engine_HP'].isnull()]

```

```

In [7]: df_temp

```

```

Out [7]:
         Make      Model  Year  Engine_Fuel_Type \
539      FIAT      500e  2015          electric
540      FIAT      500e  2016          electric
541      FIAT      500e  2017          electric
2905  Lincoln  Continental  2017  premium unleaded (recommended)
2906  Lincoln  Continental  2017  premium unleaded (recommended)
2907  Lincoln  Continental  2017  premium unleaded (recommended)
2908  Lincoln  Continental  2017  premium unleaded (recommended)
4203    Ford      Escape  2017      regular unleaded
4204    Ford      Escape  2017      regular unleaded
4205    Ford      Escape  2017      regular unleaded
4206    Ford      Escape  2017      regular unleaded
4705   Honda      Fit EV  2013          electric
4706   Honda      Fit EV  2014          electric
4785    Ford      Focus  2015          electric
4789    Ford      Focus  2016          electric
4798    Ford      Focus  2017          electric
4914    Ford  Freestar  2005      regular unleaded
4915    Ford  Freestar  2005      regular unleaded

```


4916	Ford	Freestar	2005		regular unleaded
4917	Ford	Freestar	2005		regular unleaded
4918	Ford	Freestar	2005		regular unleaded
4919	Ford	Freestar	2005		regular unleaded
5778	Mitsubishi	i-MiEV	2014		electric
5825	Chevrolet	Impala	2015	flex-fuel (unleaded/natural gas)	
5830	Chevrolet	Impala	2015	flex-fuel (unleaded/natural gas)	
5831	Chevrolet	Impala	2016	flex-fuel (unleaded/natural gas)	
5833	Chevrolet	Impala	2016	flex-fuel (unleaded/natural gas)	
5839	Chevrolet	Impala	2017	flex-fuel (unleaded/natural gas)	
5840	Chevrolet	Impala	2017	flex-fuel (unleaded/natural gas)	
6385	Nissan	Leaf	2014		electric
...
6578	Mercedes-Benz	M-Class	2015		diesel
6908	Lincoln	MKZ	2017		regular unleaded
6910	Lincoln	MKZ	2017		regular unleaded
6916	Lincoln	MKZ	2017		regular unleaded
6918	Lincoln	MKZ	2017		regular unleaded
6921	Tesla	Model S	2014		electric
6922	Tesla	Model S	2014		electric
6923	Tesla	Model S	2014		electric
6924	Tesla	Model S	2014		electric
6925	Tesla	Model S	2015		electric
6926	Tesla	Model S	2015		electric
6927	Tesla	Model S	2015		electric
6928	Tesla	Model S	2015		electric
6929	Tesla	Model S	2015		electric
6930	Tesla	Model S	2016		electric
6931	Tesla	Model S	2016		electric
6932	Tesla	Model S	2016		electric
6933	Tesla	Model S	2016		electric
6934	Tesla	Model S	2016		electric
6935	Tesla	Model S	2016		electric
6936	Tesla	Model S	2016		electric
6937	Tesla	Model S	2016		electric
6938	Tesla	Model S	2016		electric
8374	Toyota	RAV4 EV	2013		electric
8375	Toyota	RAV4 EV	2014		electric
9850	Kia	Soul EV	2015		electric
9851	Kia	Soul EV	2015		electric
9852	Kia	Soul EV	2016		electric
9853	Kia	Soul EV	2016		electric
9854	Kia	Soul EV	2016		electric

	Engine_HP	Engine_Cylinders	Transmission_Type	Driven_Wheels	\
539	NaN	0.0	DIRECT_DRIVE	front wheel drive	
540	NaN	0.0	DIRECT_DRIVE	front wheel drive	
541	NaN	0.0	DIRECT_DRIVE	front wheel drive	

2905	NaN	6.0	AUTOMATIC	all wheel drive
2906	NaN	6.0	AUTOMATIC	front wheel drive
2907	NaN	6.0	AUTOMATIC	front wheel drive
2908	NaN	6.0	AUTOMATIC	all wheel drive
4203	NaN	4.0	AUTOMATIC	front wheel drive
4204	NaN	4.0	AUTOMATIC	all wheel drive
4205	NaN	4.0	AUTOMATIC	all wheel drive
4206	NaN	4.0	AUTOMATIC	front wheel drive
4705	NaN	0.0	DIRECT_DRIVE	front wheel drive
4706	NaN	0.0	DIRECT_DRIVE	front wheel drive
4785	NaN	0.0	DIRECT_DRIVE	front wheel drive
4789	NaN	0.0	DIRECT_DRIVE	front wheel drive
4798	NaN	0.0	DIRECT_DRIVE	front wheel drive
4914	NaN	6.0	AUTOMATIC	front wheel drive
4915	NaN	6.0	AUTOMATIC	front wheel drive
4916	NaN	6.0	AUTOMATIC	front wheel drive
4917	NaN	6.0	AUTOMATIC	front wheel drive
4918	NaN	6.0	AUTOMATIC	front wheel drive
4919	NaN	6.0	AUTOMATIC	front wheel drive
5778	NaN	NaN	DIRECT_DRIVE	rear wheel drive
5825	NaN	6.0	AUTOMATIC	front wheel drive
5830	NaN	6.0	AUTOMATIC	front wheel drive
5831	NaN	6.0	AUTOMATIC	front wheel drive
5833	NaN	6.0	AUTOMATIC	front wheel drive
5839	NaN	6.0	AUTOMATIC	front wheel drive
5840	NaN	6.0	AUTOMATIC	front wheel drive
6385	NaN	0.0	DIRECT_DRIVE	front wheel drive
...
6578	NaN	4.0	AUTOMATIC	all wheel drive
6908	NaN	4.0	AUTOMATIC	front wheel drive
6910	NaN	4.0	AUTOMATIC	front wheel drive
6916	NaN	4.0	AUTOMATIC	front wheel drive
6918	NaN	4.0	AUTOMATIC	front wheel drive
6921	NaN	0.0	DIRECT_DRIVE	rear wheel drive
6922	NaN	0.0	DIRECT_DRIVE	rear wheel drive
6923	NaN	0.0	DIRECT_DRIVE	all wheel drive
6924	NaN	0.0	DIRECT_DRIVE	rear wheel drive
6925	NaN	0.0	DIRECT_DRIVE	rear wheel drive
6926	NaN	0.0	DIRECT_DRIVE	all wheel drive
6927	NaN	0.0	DIRECT_DRIVE	all wheel drive
6928	NaN	0.0	DIRECT_DRIVE	all wheel drive
6929	NaN	0.0	DIRECT_DRIVE	rear wheel drive
6930	NaN	0.0	DIRECT_DRIVE	all wheel drive
6931	NaN	0.0	DIRECT_DRIVE	all wheel drive
6932	NaN	0.0	DIRECT_DRIVE	all wheel drive
6933	NaN	0.0	DIRECT_DRIVE	rear wheel drive
6934	NaN	0.0	DIRECT_DRIVE	all wheel drive
6935	NaN	0.0	DIRECT_DRIVE	all wheel drive

6936	NaN	0.0	DIRECT_DRIVE	all wheel drive
6937	NaN	0.0	DIRECT_DRIVE	all wheel drive
6938	NaN	0.0	DIRECT_DRIVE	rear wheel drive
8374	NaN	0.0	DIRECT_DRIVE	front wheel drive
8375	NaN	0.0	DIRECT_DRIVE	front wheel drive
9850	NaN	0.0	DIRECT_DRIVE	front wheel drive
9851	NaN	0.0	DIRECT_DRIVE	front wheel drive
9852	NaN	0.0	DIRECT_DRIVE	front wheel drive
9853	NaN	0.0	DIRECT_DRIVE	front wheel drive
9854	NaN	0.0	DIRECT_DRIVE	front wheel drive

	Number_of_Doors	Vehicle_Size	Vehicle_Style	highway_MPG	city_mpg	\
539	2.0	Compact	2dr Hatchback	108	122	
540	2.0	Compact	2dr Hatchback	103	121	
541	2.0	Compact	2dr Hatchback	103	121	
2905	4.0	Large	Sedan	25	17	
2906	4.0	Large	Sedan	27	18	
2907	4.0	Large	Sedan	27	18	
2908	4.0	Large	Sedan	25	17	
4203	4.0	Compact	4dr SUV	30	23	
4204	4.0	Compact	4dr SUV	28	22	
4205	4.0	Compact	4dr SUV	28	22	
4206	4.0	Compact	4dr SUV	30	23	
4705	4.0	Compact	4dr Hatchback	105	132	
4706	4.0	Compact	4dr Hatchback	105	132	
4785	4.0	Compact	4dr Hatchback	99	110	
4789	4.0	Compact	4dr Hatchback	99	110	
4798	4.0	Compact	4dr Hatchback	99	110	
4914	4.0	Midsize	Passenger Minivan	22	16	
4915	4.0	Midsize	Passenger Minivan	22	16	
4916	4.0	Midsize	Cargo Minivan	22	16	
4917	4.0	Midsize	Passenger Minivan	22	16	
4918	4.0	Midsize	Passenger Minivan	21	16	
4919	4.0	Midsize	Passenger Minivan	21	16	
5778	4.0	Compact	4dr Hatchback	99	126	
5825	4.0	Large	Sedan	25	17	
5830	4.0	Large	Sedan	25	17	
5831	4.0	Large	Sedan	25	17	
5833	4.0	Large	Sedan	25	17	
5839	4.0	Large	Sedan	25	17	
5840	4.0	Large	Sedan	25	17	
6385	4.0	Compact	4dr Hatchback	101	126	
...	
6578	4.0	Midsize	4dr SUV	29	22	
6908	4.0	Midsize	Sedan	38	41	
6910	4.0	Midsize	Sedan	38	41	
6916	4.0	Midsize	Sedan	38	41	
6918	4.0	Midsize	Sedan	38	41	

6921	4.0	Large	Sedan	90	88
6922	4.0	Large	Sedan	97	94
6923	4.0	Large	Sedan	94	86
6924	4.0	Large	Sedan	90	88
6925	4.0	Large	Sedan	97	94
6926	4.0	Large	Sedan	102	101
6927	4.0	Large	Sedan	106	95
6928	4.0	Large	Sedan	98	89
6929	4.0	Large	Sedan	90	88
6930	NaN	Large	Sedan	105	102
6931	NaN	Large	Sedan	101	98
6932	NaN	Large	Sedan	105	92
6933	NaN	Large	Sedan	100	97
6934	NaN	Large	Sedan	107	101
6935	4.0	Large	Sedan	102	101
6936	4.0	Large	Sedan	107	101
6937	4.0	Large	Sedan	100	91
6938	4.0	Large	Sedan	90	88
8374	4.0	Midsize	4dr SUV	74	78
8375	4.0	Midsize	4dr SUV	74	78
9850	4.0	Compact	Wagon	92	120
9851	4.0	Compact	Wagon	92	120
9852	4.0	Compact	Wagon	92	120
9853	4.0	Compact	Wagon	92	120
9854	4.0	Compact	Wagon	92	120

	Popularity	MSRP
539	819	31800
540	819	31800
541	819	31800
2905	61	55915
2906	61	62915
2907	61	53915
2908	61	64915
4203	5657	29100
4204	5657	30850
4205	5657	26850
4206	5657	25100
4705	2202	36625
4706	2202	36625
4785	5657	29170
4789	5657	29170
4798	5657	29120
4914	5657	28030
4915	5657	23930
4916	5657	21630
4917	5657	26530
4918	5657	29030

4919	5657	32755
5778	436	22995
5825	1385	40660
5830	1385	37535
5831	1385	40810
5833	1385	37570
5839	1385	37675
5840	1385	40915
6385	2009	35020
...
6578	617	49800
6908	61	35010
6910	61	39510
6916	61	36760
6918	61	47670
6921	1391	79900
6922	1391	69900
6923	1391	104500
6924	1391	93400
6925	1391	69900
6926	1391	75000
6927	1391	85000
6928	1391	105000
6929	1391	80000
6930	1391	79500
6931	1391	66000
6932	1391	134500
6933	1391	74500
6934	1391	71000
6935	1391	75000
6936	1391	89500
6937	1391	112000
6938	1391	70000
8374	2031	49800
8375	2031	49800
9850	1720	35700
9851	1720	33700
9852	1720	33950
9853	1720	31950
9854	1720	35950

[69 rows x 15 columns]

There are some specific models of cars that are missing horsepower values. Let's check which models are these

```
In [8]: df_temp['Model'].unique()
```

```
Out[8]: array(['500e', 'Continental', 'Escape', 'Fit EV', 'Focus', 'Freestar',
```

```
'i-MiEV', 'Impala', 'Leaf', 'M-Class', 'MKZ', 'Model S', 'RAV4 EV',  
'Soul EV'], dtype=object)
```

One may be tempted to drop these values, but with a little help from Google, we can find the missing horsepower values and add them to the dataframe

```
In [9]: #First make a new copy of the dataframe to work with  
df = data.copy()
```

```
In [10]: #Here we add values of missing horsepower  
result = []  
for i in df['Model']:  
    if i == '500e':  
        result.append(111)  
    elif i == 'Continental':  
        result.append(400)  
    elif i == 'Escape':  
        result.append(168)  
    elif i == 'Fit EV':  
        result.append(123)  
    elif i == 'Focus':  
        result.append(143)  
    elif i == 'Freestar':  
        result.append(201)  
    elif i == 'i-MiEV':  
        result.append(66)  
    elif i == 'Impala':  
        result.append(305)  
    elif i == 'Leaf':  
        result.append(107)  
    elif i == 'M-Class':  
        result.append(201)  
    elif i == 'MKZ':  
        result.append(245)  
    elif i == 'Model S':  
        result.append(600)  
    elif i == 'RAV4 EV':  
        result.append(154)  
    elif i == 'Soul EV':  
        result.append(109)  
    else:  
        result.append(" ")  
df["Result"] = result
```

```
In [11]: #Here I will replace the missing values with a blank space, so it will be easy to merge  
df["Engine_HP"] = df["Engine_HP"].fillna('')
```

```
In [12]: df.head()
```

```

Out[12]:  Make      Model  Year      Engine_Fuel_Type  Engine_HP  \
0  BMW  1 Series M  2011  premium unleaded (required)      335
1  BMW  1 Series  2011  premium unleaded (required)      300
2  BMW  1 Series  2011  premium unleaded (required)      300
3  BMW  1 Series  2011  premium unleaded (required)      230
4  BMW  1 Series  2011  premium unleaded (required)      230

      Engine_Cylinders  Transmission_Type      Driven_Wheels  Number_of_Doors  \
0                6.0          MANUAL  rear wheel drive          2.0
1                6.0          MANUAL  rear wheel drive          2.0
2                6.0          MANUAL  rear wheel drive          2.0
3                6.0          MANUAL  rear wheel drive          2.0
4                6.0          MANUAL  rear wheel drive          2.0

      Vehicle_Size  Vehicle_Style  highway_MPG  city_mpg  Popularity  MSRP  Result
0      Compact      Coupe          26          19          3916  46135
1      Compact  Convertible          28          19          3916  40650
2      Compact      Coupe          28          20          3916  36350
3      Compact      Coupe          28          18          3916  29450
4      Compact  Convertible          28          18          3916  34500

```

Next, we will merge the new and the old horsepower columns to make a single column

```

In [13]: convert_dict = {'Result': str, 'Engine_HP': str}
         df = df.astype(convert_dict)

In [14]: df['Final_HP'] = df["Engine_HP"] + df["Result"]

In [15]: df['Final_HP'] = df['Final_HP'].astype(float)

In [16]: #Drop the old HP column
         df = df.drop(columns=['Engine_HP', 'Result'])

In [17]: #Recheck null values
         df.isnull().sum()

```

```

Out[17]: Make      0
         Model      0
         Year      0
         Engine_Fuel_Type      3
         Engine_Cylinders      30
         Transmission_Type      0
         Driven_Wheels      0
         Number_of_Doors      6
         Vehicle_Size      0
         Vehicle_Style      0
         highway_MPG      0
         city_mpg      0
         Popularity      0

```

```
MSRP          0
Final_HP      0
dtype: int64
```

Engine Cylinders

```
In [18]: #Make a temp dataframe to explore the missing values in the engine cylinders
df_temp = df[df['Engine_Cylinders'].isnull()]
```

```
In [19]: df_temp
```

```
Out[19]:
```

	Make	Model	Year	Engine_Fuel_Type	\
1983	Chevrolet	Bolt EV	2017	electric	
1984	Chevrolet	Bolt EV	2017	electric	
3716	Volkswagen	e-Golf	2015	electric	
3717	Volkswagen	e-Golf	2015	electric	
3718	Volkswagen	e-Golf	2016	electric	
3719	Volkswagen	e-Golf	2016	electric	
5778	Mitsubishi	i-MiEV	2014	electric	
5779	Mitsubishi	i-MiEV	2016	electric	
5780	Mitsubishi	i-MiEV	2017	electric	
8373	Toyota	RAV4 EV	2012	electric	
8695	Mazda	RX-7	1993	regular unleaded	
8696	Mazda	RX-7	1994	regular unleaded	
8697	Mazda	RX-7	1995	regular unleaded	
8698	Mazda	RX-8	2009	premium unleaded (required)	
8699	Mazda	RX-8	2009	premium unleaded (required)	
8700	Mazda	RX-8	2009	premium unleaded (required)	
8701	Mazda	RX-8	2009	premium unleaded (required)	
8702	Mazda	RX-8	2009	premium unleaded (required)	
8703	Mazda	RX-8	2009	premium unleaded (required)	
8704	Mazda	RX-8	2009	premium unleaded (required)	
8705	Mazda	RX-8	2010	premium unleaded (required)	
8706	Mazda	RX-8	2010	premium unleaded (required)	
8707	Mazda	RX-8	2010	premium unleaded (required)	
8708	Mazda	RX-8	2010	premium unleaded (required)	
8709	Mazda	RX-8	2010	premium unleaded (required)	
8710	Mazda	RX-8	2011	premium unleaded (required)	
8711	Mazda	RX-8	2011	premium unleaded (required)	
8712	Mazda	RX-8	2011	premium unleaded (required)	
8713	Mazda	RX-8	2011	premium unleaded (required)	
8714	Mazda	RX-8	2011	premium unleaded (required)	

	Engine_Cylinders	Transmission_Type	Driven_Wheels	Number_of_Doors	\
1983	NaN	DIRECT_DRIVE	front wheel drive	4.0	
1984	NaN	DIRECT_DRIVE	front wheel drive	4.0	
3716	NaN	DIRECT_DRIVE	front wheel drive	4.0	
3717	NaN	DIRECT_DRIVE	front wheel drive	4.0	

3718	NaN	DIRECT_DRIVE	front wheel drive	4.0
3719	NaN	DIRECT_DRIVE	front wheel drive	4.0
5778	NaN	DIRECT_DRIVE	rear wheel drive	4.0
5779	NaN	DIRECT_DRIVE	rear wheel drive	4.0
5780	NaN	DIRECT_DRIVE	rear wheel drive	4.0
8373	NaN	DIRECT_DRIVE	front wheel drive	4.0
8695	NaN	MANUAL	rear wheel drive	2.0
8696	NaN	MANUAL	rear wheel drive	2.0
8697	NaN	MANUAL	rear wheel drive	2.0
8698	NaN	MANUAL	rear wheel drive	4.0
8699	NaN	AUTOMATIC	rear wheel drive	4.0
8700	NaN	MANUAL	rear wheel drive	4.0
8701	NaN	MANUAL	rear wheel drive	4.0
8702	NaN	MANUAL	rear wheel drive	4.0
8703	NaN	AUTOMATIC	rear wheel drive	4.0
8704	NaN	AUTOMATIC	rear wheel drive	4.0
8705	NaN	MANUAL	rear wheel drive	4.0
8706	NaN	AUTOMATIC	rear wheel drive	4.0
8707	NaN	AUTOMATIC	rear wheel drive	4.0
8708	NaN	MANUAL	rear wheel drive	4.0
8709	NaN	MANUAL	rear wheel drive	4.0
8710	NaN	AUTOMATIC	rear wheel drive	4.0
8711	NaN	MANUAL	rear wheel drive	4.0
8712	NaN	MANUAL	rear wheel drive	4.0
8713	NaN	MANUAL	rear wheel drive	4.0
8714	NaN	AUTOMATIC	rear wheel drive	4.0

	Vehicle_Size	Vehicle_Style	highway_MPG	city_mpg	Popularity	MSRP \
1983	Compact	4dr Hatchback	110	128	1385	40905
1984	Compact	4dr Hatchback	110	128	1385	36620
3716	Compact	4dr Hatchback	105	126	873	33450
3717	Compact	4dr Hatchback	105	126	873	35445
3718	Compact	4dr Hatchback	105	126	873	28995
3719	Compact	4dr Hatchback	105	126	873	35595
5778	Compact	4dr Hatchback	99	126	436	22995
5779	Compact	4dr Hatchback	99	126	436	22995
5780	Compact	4dr Hatchback	102	121	436	22995
8373	Midsize	4dr SUV	74	78	2031	49800
8695	Compact	Coupe	23	15	586	7523
8696	Compact	Coupe	23	15	586	8147
8697	Compact	Coupe	23	15	586	8839
8698	Compact	Coupe	22	16	586	31930
8699	Compact	Coupe	23	16	586	26435
8700	Compact	Coupe	22	16	586	27860
8701	Compact	Coupe	22	16	586	31000
8702	Compact	Coupe	22	16	586	26435
8703	Compact	Coupe	23	16	586	31700
8704	Compact	Coupe	23	16	586	28560

8705	Compact	Coupe	22	16	586	32140
8706	Compact	Coupe	23	16	586	26645
8707	Compact	Coupe	23	16	586	32810
8708	Compact	Coupe	22	16	586	26645
8709	Compact	Coupe	22	16	586	32110
8710	Compact	Coupe	23	16	586	32960
8711	Compact	Coupe	22	16	586	32260
8712	Compact	Coupe	22	16	586	32290
8713	Compact	Coupe	22	16	586	26795
8714	Compact	Coupe	23	16	586	26795

	Final_HP
1983	200.0000
1984	200.0000
3716	115.0000
3717	115.0000
3718	115.0000
3719	115.0000
5778	66.0000
5779	66.0660
5780	66.0660
8373	154.0154
8695	255.0000
8696	255.0000
8697	255.0000
8698	232.0000
8699	212.0000
8700	232.0000
8701	232.0000
8702	232.0000
8703	212.0000
8704	212.0000
8705	232.0000
8706	212.0000
8707	212.0000
8708	232.0000
8709	232.0000
8710	212.0000
8711	232.0000
8712	232.0000
8713	232.0000
8714	212.0000

These cars are either electric or Mazdas with a rotary engine. In either case, they did not have any cylinders, so we can safely replace the NaN with 0.

```
In [20]: #Replace NaN with 0 in the Engine Cylinders Column
df['Engine_Cylinders'].fillna(0, inplace = True);
```

```
In [21]: df.isnull().sum()
```

```
Out[21]: Make          0
         Model         0
         Year          0
         Engine_Fuel_Type  3
         Engine_Cylinders  0
         Transmission_Type  0
         Driven_Wheels    0
         Number_of_Doors  6
         Vehicle_Size     0
         Vehicle_Style    0
         highway_MPG      0
         city_mpg         0
         Popularity       0
         MSRP             0
         Final_HP        0
         dtype: int64
```

Engine Fuel

```
In [22]: df_temp = df[df['Engine_Fuel_Type'].isnull()]
```

```
In [23]: df_temp
```

```
Out[23]:
```

	Make	Model	Year	Engine_Fuel_Type	Engine_Cylinders	\
11321	Suzuki	Verona	2004	NaN	6.0	
11322	Suzuki	Verona	2004	NaN	6.0	
11323	Suzuki	Verona	2004	NaN	6.0	

	Transmission_Type	Driven_Wheels	Number_of_Doors	Vehicle_Size	\
11321	AUTOMATIC	front wheel drive	4.0	Midsize	
11322	AUTOMATIC	front wheel drive	4.0	Midsize	
11323	AUTOMATIC	front wheel drive	4.0	Midsize	

	Vehicle_Style	highway_MPG	city_mpg	Popularity	MSRP	Final_HP
11321	Sedan	25	17	481	17199	155.0
11322	Sedan	25	17	481	20199	155.0
11323	Sedan	25	17	481	18499	155.0

```
In [24]: #Replace NaN with regular unleaded in the Engine Fuel Type Column
         df['Engine_Fuel_Type'].fillna('regular unleaded', inplace = True);
```

Number of doors

```
In [25]: df_temp = df[df['Number_of_Doors'].isnull()]
```

```
In [26]: df_temp
```

```

Out [26]:
      Make      Model  Year      Engine_Fuel_Type  Engine_Cylinders  \
4666  Ferrari      FF  2013  premium unleaded (required)          12.0
6930   Tesla  Model S  2016                electric              0.0
6931   Tesla  Model S  2016                electric              0.0
6932   Tesla  Model S  2016                electric              0.0
6933   Tesla  Model S  2016                electric              0.0
6934   Tesla  Model S  2016                electric              0.0

      Transmission_Type  Driven_Wheels  Number_of_Doors  Vehicle_Size  \
4666  AUTOMATED_MANUAL  all wheel drive              NaN          Large
6930   DIRECT_DRIVE  all wheel drive              NaN          Large
6931   DIRECT_DRIVE  all wheel drive              NaN          Large
6932   DIRECT_DRIVE  all wheel drive              NaN          Large
6933   DIRECT_DRIVE  rear wheel drive              NaN          Large
6934   DIRECT_DRIVE  all wheel drive              NaN          Large

      Vehicle_Style  highway_MPG  city_mpg  Popularity  MSRP  Final_HP
4666          Coupe           16        11         2774  295000    651.0
6930          Sedan           105       102         1391   79500    600.0
6931          Sedan           101        98         1391   66000    600.0
6932          Sedan           105        92         1391  134500    600.0
6933          Sedan           100        97         1391   74500    600.0
6934          Sedan           107       101         1391   71000    600.0

```

In this case we see that there are two car models that don't have the number of doors. For all the Teslas we will have 4 doors, while the ferrari has 2 doors. First I will change the NaN to 4, and then deal with the ferrari later if needed.

```

In [27]: #Replace NaN with 4 in the Number of doors Column
df['Number_of_Doors'].fillna(4, inplace = True);

```

```

In [28]: df.isnull().sum()

```

```

Out [28]: Make      0
Model      0
Year       0
Engine_Fuel_Type  0
Engine_Cylinders  0
Transmission_Type  0
Driven_Wheels    0
Number_of_Doors  0
Vehicle_Size     0
Vehicle_Style    0
highway_MPG      0
city_mpg        0
Popularity       0
MSRP             0
Final_HP        0
dtype: int64

```

As we can see above, we have taken care of the null values
Next up, we will look at duplicated values

1.2.2 Duplicated Values

```
In [29]: dup_rows = df[df.duplicated()]
```

```
In [30]: dup_rows.head(20)
```

```
Out [30]:
```

	Make	Model	Year	Engine_Fuel_Type	\
14	BMW	1 Series	2013	premium unleaded (required)	
18	Audi	100	1992	regular unleaded	
20	Audi	100	1992	regular unleaded	
24	Audi	100	1993	regular unleaded	
25	Audi	100	1993	regular unleaded	
88	Nissan	200SX	1996	regular unleaded	
92	Nissan	200SX	1997	regular unleaded	
94	Nissan	200SX	1998	regular unleaded	
109	Volvo	240	1992	regular unleaded	
126	BMW	3 Series Gran Turismo	2015	premium unleaded (required)	
137	BMW	3 Series	2015	premium unleaded (required)	
141	BMW	3 Series	2015	premium unleaded (required)	
252	Mazda	323	1992	regular unleaded	
413	BMW	4 Series Gran Coupe	2015	premium unleaded (required)	
414	BMW	4 Series Gran Coupe	2015	premium unleaded (required)	
431	BMW	4 Series	2015	premium unleaded (required)	
432	BMW	4 Series	2015	premium unleaded (required)	
435	BMW	4 Series	2015	premium unleaded (required)	
436	BMW	4 Series	2015	premium unleaded (required)	
677	Pontiac	6000	1990	regular unleaded	

	Engine_Cylinders	Transmission_Type	Driven_Wheels	Number_of_Doors	\
14	6.0	MANUAL	rear wheel drive	2.0	
18	6.0	MANUAL	front wheel drive	4.0	
20	6.0	MANUAL	front wheel drive	4.0	
24	6.0	MANUAL	front wheel drive	4.0	
25	6.0	MANUAL	front wheel drive	4.0	
88	4.0	MANUAL	front wheel drive	2.0	
92	4.0	MANUAL	front wheel drive	2.0	
94	4.0	MANUAL	front wheel drive	2.0	
109	4.0	MANUAL	rear wheel drive	4.0	
126	4.0	AUTOMATIC	all wheel drive	4.0	
137	4.0	AUTOMATIC	all wheel drive	4.0	
141	4.0	AUTOMATIC	rear wheel drive	4.0	
252	4.0	MANUAL	front wheel drive	2.0	
413	4.0	AUTOMATIC	rear wheel drive	4.0	
414	4.0	AUTOMATIC	all wheel drive	4.0	
431	4.0	AUTOMATIC	all wheel drive	2.0	

432	4.0	AUTOMATIC	all wheel drive	2.0
435	4.0	AUTOMATIC	rear wheel drive	2.0
436	4.0	AUTOMATIC	rear wheel drive	2.0
677	6.0	AUTOMATIC	front wheel drive	4.0


	Vehicle_Size	Vehicle_Style	highway_MPG	city_mpg	Popularity	MSRP	\
14	Compact	Coupe	28	19	3916	31500	
18	Midsize	Sedan	24	17	3105	2000	
20	Midsize	Sedan	24	17	3105	2000	
24	Midsize	Sedan	24	17	3105	2000	
25	Midsize	Sedan	24	17	3105	2000	
88	Compact	Coupe	36	26	2009	2000	
92	Compact	Coupe	35	25	2009	2000	
94	Compact	Coupe	35	25	2009	2000	
109	Midsize	Sedan	26	19	870	2000	
126	Midsize	4dr Hatchback	33	22	3916	41850	
137	Midsize	Sedan	33	22	3916	39500	
141	Midsize	Sedan	35	23	3916	37500	
252	Compact	2dr Hatchback	33	25	586	2000	
413	Midsize	Sedan	34	23	3916	40300	
414	Midsize	Sedan	33	22	3916	42300	
431	Midsize	Coupe	33	22	3916	42750	
432	Midsize	Convertible	33	21	3916	50750	
435	Midsize	Convertible	34	23	3916	48750	
436	Midsize	Coupe	35	23	3916	40750	
677	Midsize	Wagon	27	17	210	2000	

	Final_HP
14	230.0
18	172.0
20	172.0
24	172.0
25	172.0
88	115.0
92	115.0
94	115.0
109	114.0
126	240.0
137	240.0
141	240.0
252	82.0
413	240.0
414	240.0
431	240.0
432	240.0
435	240.0
436	240.0
677	140.0

We see that there are duplicated values, but it just could be multiple cars with the same attributes, and therefore, I will not drop them

1.2.3 Data exploration for outliers

```
In [31]: df.hist(bins = 20,grid=True, figsize = (12,12), color = 'red');  
plt.tight_layout()
```



Observations

- There are some outliers in the price. Even though it is true that there are cars that are USD 2M, they will completely skew the distribution, so we will drop any car >500K
- We will also limit highway MPG to 130.

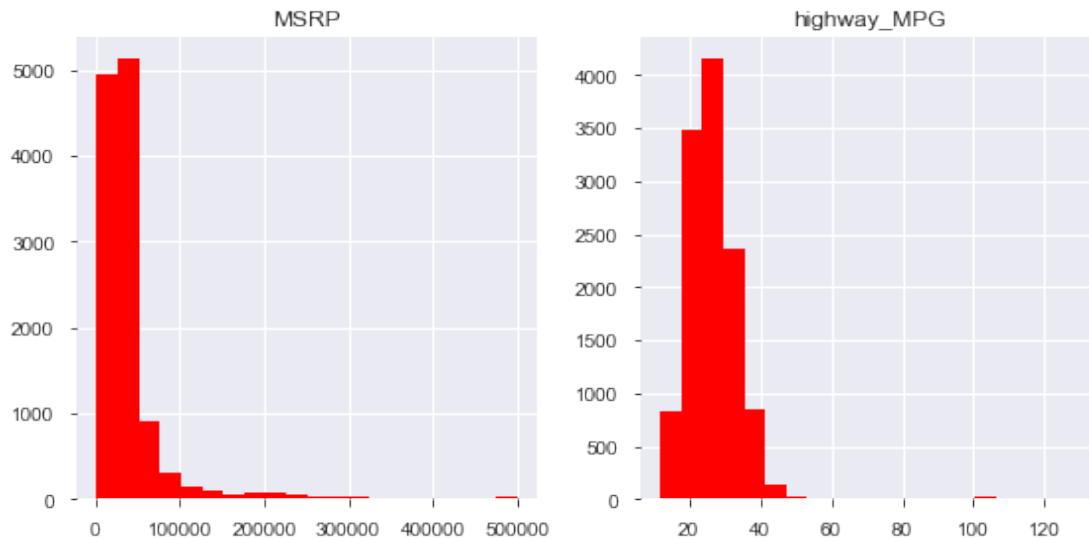
Limit the values of price and highway MPG

```
In [32]: df_clean = df.copy()
```

```
In [33]: df_clean['MSRP'] = df_clean['MSRP'].clip(upper = 500000)
```

```
In [34]: df_clean['highway_MPG'] = df_clean['highway_MPG'].clip(upper = 130)
```

```
In [35]: df_clean.hist(column = ['highway_MPG', 'MSRP'], bins = 20, grid=True, figsize = (8,4), c
plt.tight_layout()
```




We can see that the data is still right skewed, and we will log transform this before applying ML algorithms

Now that we have cleaned up the data, we will generate a profile report


1.3 Univariate Data Exploration

```
In [36]: #We will write a function that we can use for univariate exploration
#Variable is the column name
#x and y are length and width of the plot
def univar(variable, x, y):
    plt.figure(figsize=[x, y])
    sns.countplot(data = df_clean, x= variable, order = df_clean[variable].value_counts().index)
    plt.xticks(rotation=90);


In [37]: #Make of the car
univar('Make', 14, 5)
```


```
In [38]: #Engine fuel type
         univar('Engine_Fuel_Type', 5, 5)
```




```
In [39]: #Vehicle Size
         univar('Vehicle_Size', 5, 5)
```




```
In [40]: #Vehicle Style  
         univar('Vehicle_Style',8,5)
```



```
In [41]: #Transmission
         univar('Transmission_Type',5,5)
```



```
In [42]: #Drive wheels
         univar('Driven_Wheels',5,5)
```




1.3.1 Univariate Data Exploration Summary

Most common categories - Brand: Chevy - Fuel: Regular unleaded - Type: Sedan - Transmission: Automatic - Drive: Front Wheel Drive

1.4 Bivariate Data Exploration

```
In [43]: # correlation plot
plt.figure(figsize = [12,12])
sns.heatmap(df_clean.corr(), annot = True, fmt = '.2f',
            cmap = 'vlag_r', center = 0)
plt.show()
```




1.4.1 MSRP

In [44]: *#MSRP vs Transmission and Wheels driven*


```
fig, axarr = plt.subplots(1, 2, figsize=(12, 6), sharey = True)
```

```
sns.barplot(x = 'Transmission_Type', y = 'MSRP', data = df_clean,palette="Reds", ax=axarr[0])
sns.barplot(x = 'Driven_Wheels', y = 'MSRP', data = df_clean,palette="Reds", ax=axarr[1])
plt.sca(axarr[0])
plt.xticks(rotation=90)
plt.sca(axarr[1])
plt.xticks(rotation=90)
plt.tight_layout()
```




In [45]: *#Let's check the effect of vehicle size on the MSRP*

```
plt.figure(figsize=[8, 5])
sns.barplot(x = 'Vehicle_Style', y = 'MSRP', data = df_clean,palette="Reds");
plt.xticks(rotation=90);
```



In [46]: *#Car Make vs MSRP*

```
plt.figure(figsize=[14, 5])
sns.barplot(x = 'Make', y = 'MSRP', data = df_clean,palette="Reds");
plt.yscale('symlog')
plt.ylim(15000,500000)
plt.xticks(rotation=90);
```




1.4.2 Highway MPG


In [47]: *#MSRP vs Transmission and Wheels driven*

```
fig, axarr = plt.subplots(1, 2, figsize=(12, 6), sharey = True)
```


```
sns.barplot(x = 'Transmission_Type', y = 'highway_MPG', data = df_clean,palette="Greens")
sns.barplot(x = 'Driven_Wheels', y = 'highway_MPG', data = df_clean,palette="Greens",
plt.sca(axarr[0])
plt.xticks(rotation=90)
plt.sca(axarr[1])
plt.xticks(rotation=90)
plt.tight_layout()
```



```
In [48]: #Let's check the effect of vehicle size on the Mileage
plt.figure(figsize=[8, 5])
sns.barplot(x = 'Vehicle_Style', y = 'highway_MPG', data = df_clean,palette="Greens")
plt.xticks(rotation=90);
```



```
In [49]: #Car Make vs MPG
plt.figure(figsize=[14, 5])
sns.barplot(x = 'Make', y = 'highway_MPG', data = df_clean,palette="Greens");
#plt.ylim(15000,500000)
plt.xticks(rotation=90);
```



1.4.3 Bivariate Data Exploration Summary

MSRP

- Automated Manuals are most expensive (these are mostly exotic supercars)
- FWD are cheapest
- Coupes and Convertibles are most expensive
- There are three tiers based on brands
 - Tier 1 consists of ultra premium cars are Bugatti, Maybach, Ferrari, etc
 - The second tier is BMW, Audi, Mercedes, Infiniti etc
 - Most mass market cars like Ford, Chevy form the third tier

MPG

- Direct drive has highest MPG (these electric cars)
- FWD cars have better MPG than rear wheel or AWD (these are more mass market cars, with less performance)
- 4DR Hatchbacks have best MPG (most electric cars fall in this category)
- Tesla as a brand has the best MPG (They only make electric cars)


1.5 Multivariate Exploration

In [50]: *#Let's first look at the MSRP vs Horsepower*


```
plt.figure(figsize=[9, 6])
sns.set_style("darkgrid")
sns.scatterplot(data = df_clean, x = 'Final_HP', y = 'MSRP', hue = 'Engine_Cylinders')

#Rescale the plot to better visualize the distribution
plt.xscale('log')
plt.xticks([50, 100, 200, 500, 1000], [50, 100, 200, 500, 1000])
```

```
plt.yscale('log')
plt.yticks([ 5000,20000, 100000, 500000], [ '5k' , '20k', '100k', '500k']);
```



```
In [51]: plt.figure(figsize=[9, 6])
sns.set_style("darkgrid")
sns.scatterplot(data = df_clean, x = 'Final_HP', y = 'highway_MPG', hue = 'Engine_Cyl')
plt.xscale('log')
plt.xticks([50, 100, 200, 500, 1000], [50, 100, 200, 500, 1000])
plt.yscale('log')
plt.yticks([20, 50, 100], [ 20, 50, 100]);
```



1.6 Building a classification model for Fuel Economy

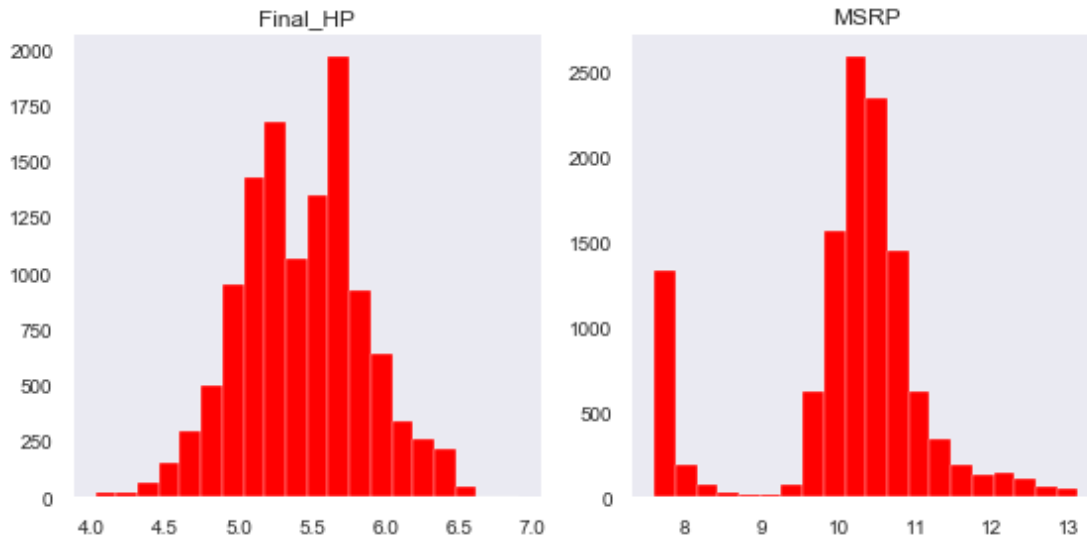
```
In [83]: df_class = df_clean.copy()
```

1.6.1 Transform and scale the dataset

```
In [84]: # Split the data into features and target label
mileage_raw = df_class[['highway_MPG']]
# in the features dataset, we will drop the highway_MPG, because that is the target
# We will also drop city_mpg because that is highly correlated to highway_MPG,
# and Engine_Cylinders, which are highly correlated to Horsepower
features_raw = df_class.drop(['highway_MPG', 'city_mpg', 'Engine_Cylinders'], axis = 1)

In [85]: # Log-transform the skewed features
skewed = ['MSRP', 'Final_HP']
features_log_transformed = pd.DataFrame(data = features_raw)
features_log_transformed[skewed] = features_raw[skewed].apply(lambda x: np.log(x + 1))

In [86]: features_log_transformed.hist(column = ['Final_HP', 'MSRP'], bins = 20, grid=False, figs:
plt.tight_layout()
```



```
In [88]: # Import sklearn.preprocessing.StandardScaler
from sklearn.preprocessing import MinMaxScaler

# Initialize a scaler, then apply it to the features
scaler = MinMaxScaler() # default=(0, 1)
numerical = ['Year', 'Number_of_Doors', 'Final_HP']

features_log_minmax_transform = pd.DataFrame(data = features_log_transformed)
features_log_minmax_transform[numerical] = scaler.fit_transform(features_log_transformed[numerical])

# Show an example of a record with scaling applied
display(features_log_minmax_transform.head(n = 10))
```

	Make	Model	Year	Engine_Fuel_Type	Transmission_Type	\
0	BMW	1 Series M	0.777778	premium unleaded (required)	MANUAL	
1	BMW	1 Series	0.777778	premium unleaded (required)	MANUAL	
2	BMW	1 Series	0.777778	premium unleaded (required)	MANUAL	
3	BMW	1 Series	0.777778	premium unleaded (required)	MANUAL	
4	BMW	1 Series	0.777778	premium unleaded (required)	MANUAL	
5	BMW	1 Series	0.814815	premium unleaded (required)	MANUAL	
6	BMW	1 Series	0.814815	premium unleaded (required)	MANUAL	
7	BMW	1 Series	0.814815	premium unleaded (required)	MANUAL	
8	BMW	1 Series	0.814815	premium unleaded (required)	MANUAL	
9	BMW	1 Series	0.851852	premium unleaded (required)	MANUAL	

	Driven_Wheels	Number_of_Doors	Vehicle_Size	Vehicle_Style	Popularity	\
0	rear wheel drive	0.0	Compact	Coupe	3916	
1	rear wheel drive	0.0	Compact	Convertible	3916	
2	rear wheel drive	0.0	Compact	Coupe	3916	

3	rear wheel drive	0.0	Compact	Coupe	3916
4	rear wheel drive	0.0	Compact	Convertible	3916
5	rear wheel drive	0.0	Compact	Coupe	3916
6	rear wheel drive	0.0	Compact	Convertible	3916
7	rear wheel drive	0.0	Compact	Coupe	3916
8	rear wheel drive	0.0	Compact	Convertible	3916
9	rear wheel drive	0.0	Compact	Convertible	3916


	MSRP	Final_HP
0	10.739349	0.621189
1	10.612779	0.583053
2	10.500977	0.583053
3	10.290483	0.491286
4	10.448744	0.491286
5	10.348205	0.491286
6	10.694238	0.583053
7	10.579005	0.583053
8	10.515994	0.491286
9	10.524091	0.491286

```
In [89]: mileage_raw.head()
```

```
Out[89]:
```

	highway_MPG
0	26
1	28
2	28
3	28
4	28

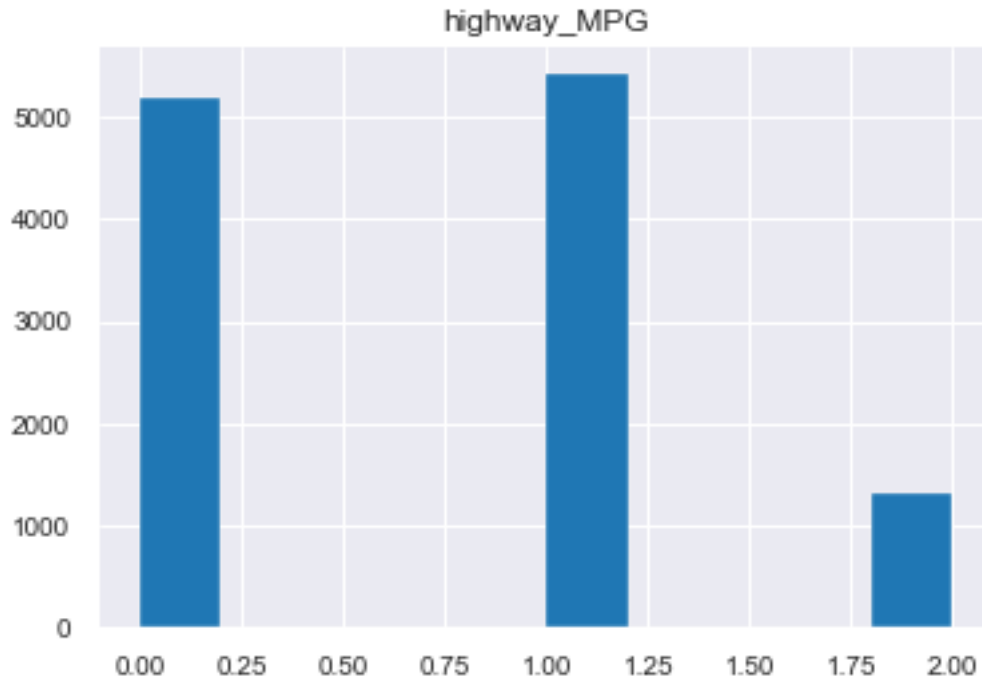
```
In [90]: mileage_raw.hist();
```

```
In [91]: #We will split the mileage (MPG) into three categories
mileage_raw.loc[(mileage_raw['highway_MPG'] < 25), 'highway_MPG'] = 0
mileage_raw.loc[(mileage_raw['highway_MPG'] >= 25)&(mileage_raw['highway_MPG'] < 35),
mileage_raw.loc[(mileage_raw['highway_MPG'] >= 35), 'highway_MPG'] = 2
```

```
In [92]: mileage_raw.hist()
```

```
Out[92]: array([[<matplotlib.axes._subplots.AxesSubplot object at 0x1a1e1d2e10>]],
dtype=object)
```



```
In [93]: # One-hot encode the 'features_log_minmax_transform' data using pandas.get_dummies()
features_final = pd.get_dummies(features_log_minmax_transform)
```

```
# Print the number of features after one-hot encoding
encoded = list(features_final.columns)
print("{} total features after one-hot encoding.".format(len(encoded)))

print (encoded)
```

1006 total features after one-hot encoding.

['Year', 'Number_of_Doors', 'Popularity', 'MSRP', 'Final_HP', 'Make_Acura', 'Make_Alfa Romeo',

```
In [94]: # Import train_test_split
from sklearn.model_selection import train_test_split

# Split the 'features' and 'income' data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(features_final,
                                                    mileage_raw,
                                                    test_size = 0.2,
                                                    random_state = 0)

# Show the results of the split
print("Training set has {} samples.".format(X_train.shape[0]))
print("Testing set has {} samples.".format(X_test.shape[0]))
```

Training set has 9531 samples.
Testing set has 2383 samples.

```
In [95]: #import necessary ML libraries
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier, AdaBoostClassifier, GradientBoos
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
from sklearn.metrics import make_scorer, confusion_matrix, fbeta_score

In [96]: #Pick three classifiers
clf_A = LogisticRegression(random_state = 0)
clf_B = GradientBoostingClassifier(random_state = 0)
clf_C = RandomForestClassifier(random_state = 0)

#Fit the data to the three classifiers and print results from them
for clf in [clf_A, clf_B, clf_C]:
    learner = clf
    start = time() # Get start time
    learner = learner.fit(X_train, y_train)
    end = time() # Get start time
    train_time = end-start

    predictions_test = learner.predict(X_test)
    predictions_train = learner.predict(X_train)

    # Score our model
    print(clf.__class__.__name__)
    print('Accuracy score: ', format(accuracy_score(y_test, predictions_test)))
    print('Precision score: ', format(precision_score(y_test, predictions_test,average='wei
    print('Recall score: ', format(recall_score(y_test, predictions_test,average='wei
    print('F1 score: ', format(f1_score(y_test, predictions_test,average='weighted')))
    print('Time: ', format(train_time))

    print('Confusion Matrix')
    print(confusion_matrix(y_test,predictions_test))
    print( )

LogisticRegression
Accuracy score:  0.9148132605958875
Precision score:  0.9145730678269199
Recall score:    0.9148132605958875
F1 score:        0.9144162757920861
Time: 0.36118006706237793
Confusion Matrix
[[ 969   60    0]
```

```
[ 64 1007  27]
[   0   52 204]]
```

GradientBoostingClassifier

```
Accuracy score:  0.9093579521611415
Precision score:  0.9103244617115024
Recall score:    0.9093579521611415
F1 score:        0.9092466847423099
Time: 41.488425970077515
Confusion Matrix
[[ 941   88    0]
 [  54 1019   25]
 [   1   48 207]]
```

RandomForestClassifier

```
Accuracy score:  0.9618128409567772
Precision score:  0.9617751299666105
Recall score:    0.9618128409567772
F1 score:        0.9617689108879187
Time: 0.35206103324890137
Confusion Matrix
[[1005   24    0]
 [  36 1048   14]
 [   0   17 239]]
```

```
In [127]: from sklearn.model_selection import GridSearchCV
```

```
#Do a GridSearch to optimize the best model from the previous step
```

```
# build a classifier
```

```
clf_rf = RandomForestClassifier(random_state = 0)
```

```
# Set up the hyperparameter search
```

```
parameters = {"n_estimators": [10,50,100] , "max_depth": [5,50,250], "min_samples_split": [2,5,10]}
```

```
# Run a randomized search over the hyperparameters
```

```
random_search = GridSearchCV(clf_rf, parameters)
```

```
# Fit the model on the training data
```

```
grid_fit = random_search.fit(X_train, y_train)
```

```
#Get the estimator
```

```
best_clf = grid_fit.best_estimator_
```

```
# Make predictions on the test data
```

```
#rf_preds = random_search.best_estimator_.predict(X_test)
```

```

predictions = (clf.fit(X_train, y_train)).predict(X_test)
best_predictions = best_clf.predict(X_test)

print('Accuracy score: ', format(accuracy_score(y_test, predictions)))
print('Precision score: ', format(precision_score(y_test, predictions, average='weighted')))
print('Recall score: ', format(recall_score(y_test, predictions, average='weighted')))
print('F1 score: ', format(f1_score(y_test, predictions, average='weighted')))
print('\n\n')
print('Confusion Matrix')
print(confusion_matrix(y_test, predictions))

```

```

Accuracy score:  0.9618128409567772
Precision score:  0.9617751299666105
Recall score:    0.9618128409567772
F1 score:        0.9617689108879187

```

```

Confusion Matrix
[[1005   24    0]
 [  36 1048   14]
 [    0   17  239]]

```

```

In [128]: #Make a function to show the top features
def feature_plot(importances, X_train, y_train):

    # Display the five most important features
    indices = np.argsort(importances)[::-1]
    columns = X_train.columns.values[indices[:8]]
    values = importances[indices][:8]

    # Creat the plot
    fig = plt.figure(figsize = (12,7))
    plt.title("Normalized Weights for First Five Most Predictive Features", fontsize=16)
    plt.bar(np.arange(8), values, width = 0.6, align="center", color = '#00A000', \
            label = "Feature Weight")
    plt.bar(np.arange(8) - 0.3, np.cumsum(values), width = 0.2, align = "center", color = '#FF0000', \
            label = "Cumulative Feature Weight")
    plt.xticks(np.arange(8), columns)
    plt.xlim((-0.5, 7.5))
    plt.ylabel("Weight", fontsize = 12)
    plt.xlabel("Feature", fontsize = 12)
    plt.xticks(rotation=45);
    plt.legend(loc = 'upper center')
    plt.tight_layout()
    plt.show()

```

```

In [129]: # Import a supervised learning model that has 'feature_importances_'

```

```

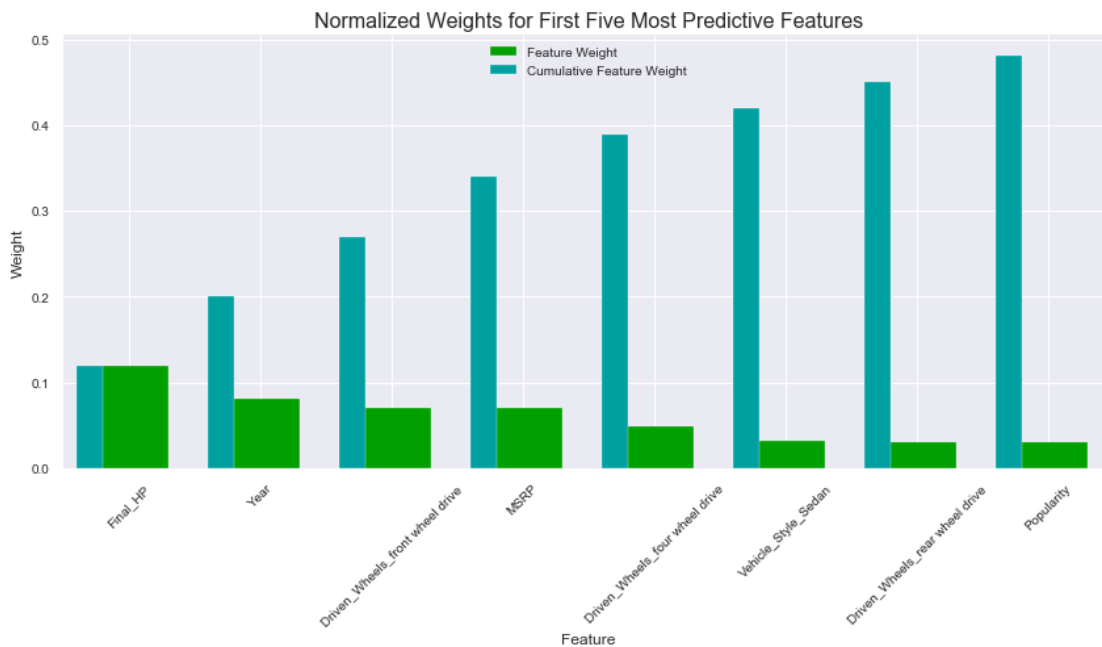
###Importing a new model is not necessary as the GradientBoostingClassifier already i

# Train the supervised model on the training set using .fit(X_train, y_train)
model = best_clf

# Extract the feature importances using .feature_importances_
importances = model.feature_importances_

# Plot
feature_plot(importances, X_train, y_train)

```



We see that the Final_HP is the top predictive feature. This is similar to what we found in our previous bivariate analysis, where we saw that HP and MPG were inversely related

In []: